

Paritosh Gandre, M.S.

Mountain View, CA | +1 330-554-8445 | paritoshkrcg@gmail.com | github.com/paritosh100 |
linkedin.com/in/paritosh-gandre

Professional Summary

Full-stack engineer with 3 years shipping production-grade applications end to end, from API design through frontend. Deep understanding of system architecture demonstrated by scalable, high-volume backend systems serving tens of thousands of concurrent users at 99.9% SLA and sub-1000ms latency. Owns architecture decisions from multi-tenant security to AI/LLM integration, raising the reliability bar on production AI features.

Technical Skills

Frontend: React 19, Next.js 16, TypeScript, Zod

API Design & Backend: Python (FastAPI, Flask), REST API design, async orchestration, Pydantic-validated contracts

Cloud & DevOps: AWS (EC2, NLB, RDS, IAM), Docker, Terraform, Google Cloud Run, Vercel, CI/CD (GitHub Actions, cloudbuild.yaml)

System Architecture & Databases: PostgreSQL, Row-Level Security (RLS), advisory-lock transactions, ACID-compliant schema design, Supabase

AI & LLM Integration: PyTorch, TensorFlow, RAG, FAISS, LLM-as-judge evaluation, multi-agent orchestration, dual-LLM backend routing

Reliability & Compliance: Datadog observability, automated testing (unit/integration/e2e), GDPR-compliant audit design, TOTP-based MFA, RBAC

Professional Experience

Senior Machine Learning Engineer

Nov 2025 – Present

OViiE AI (NVIDIA Inception) | Remote

- Architected a production-grade, end-to-end ML inference system from first principles, integrating cloud speech-to-text, custom translation model deployment, and real-time speech synthesis over WebRTC, maintaining sub-1000ms latency in production.
- Designed scalable, high-volume system architecture for a multi-tenant platform serving dozens of hospitality properties; shipped 6 production systems end-to-end.
- Implemented PostgreSQL Row-Level Security (RLS) policies to enforce Guest/Agent/Admin access boundaries with zero cross-org incidents; built GDPR-compliant audit trails and TOTP-based multi-factor authentication; designed role-based access control (RBAC) across a service-role write architecture.
- Achieved 90%+ test coverage across database, API, and frontend layers with 2,000+ automated tests, maintaining zero customer-facing billing errors and zero production errors in a live multi-tenant environment.
- Built an AI-graded multi-model evaluation harness using LLM-as-judge methodology to identify and resolve failure modes, informing model migration; reduced TTS synthesis costs 80% through caching optimization, delivering multi-five-figure annual infrastructure savings.

Data Scientist / Machine Learning Engineer

May 2024 – Dec 2024

Inke (SafeSpace) | Remote

- Built NLP and LLM pipelines handling 100K+ daily requests at 85% accuracy and sub-200ms p99 latency.
- Reduced inference latency 30% through request batching, model quantization, and response caching.

Data Scientist

Jan 2022 – Jul 2023

SL Kulkarni | Mumbai, India

- Deployed anomaly detection across 500+ real-time sensor streams, achieving 95% precision and 85% recall.
- Cut equipment downtime 20% and saved \$50K+ annually in prevented maintenance costs.

Projects

InsightBridge - Multi-Agent LLM Orchestration Platform

github.com/paritosh100/InsightBridge

- Designed a 6-agent orchestration pipeline using API design principles on FastAPI and Next.js, reducing LLM output errors through async orchestration and Pydantic-validated contracts.
- Optimized LLM response latency 4x (2–3s vs. 8–12s) via feature-flagged backend routing, with content-safety guardrails codified across dual LLM backends (OpenAI + Gemini).
- Deployed a two-service production application (Vercel frontend, Docker/Cloud Run backend) with DevOps-ready CI configs enabling independent scaling.

Education

M.S., Data Science

Aug 2023 – May 2025

Kent State University | Kent, OH

GPA: 3.4/4.0