

# Paritosh Kiran Gandre | DATA SCIENTIST

+1 330 554 8445 | paritoshkrcg@gmail.com | Portfolio | Github | LinkedIn

## Experience

### AI Agents Engineer (Industry Project) - Remote

Nov 2025 – Dec 2025

Google x Kaggle AI Agents Intensive

- Built MediGuard AI, a multi-agent healthcare fraud detection system analyzing Synthea patient claims (ICD-10, CPT, HCPCS) using Google ADK and Gemini LLMs
- Designed 3 specialized LLM agents for identity fraud, billing anomaly, and discharge readiness analysis using SequentialAgent orchestration
- Developed 7 FunctionTools in Python for claims ETL, demographics consistency checks, ICD-10 and CPT validation, anomaly detection, and statistical cost analysis using Pandas
- Built a React and Next.js dashboard with progressive multi-agent workflow visualization supporting healthcare audit teams

### Research Assistant – Deep Learning for Genomic Data Analysis

Feb 2025 – May 2025

Kent State University, Kent, OH

- Drove 97% CNN accuracy on 40K DNA sequences by curating datasets and tuning convolutional blocks in TensorFlow/PyTorch, validated with ROC-AUC and F1
- Cut manual data handling by 80% by building Dockerized ETL pipelines (pandas/Dask), versioned data with DVC, and tracked runs in MLflow

### Data Analyst – Supply Chain

Jan 2025 – May 2025

Kent State University, Kent, OH

- Reduced over-purchasing 15% (\$10K) by shipping ARIMA/SARIMAX demand models on 3 years of POS & inventory and monitoring MAPE and bias
- Cut report latency 30% by automating SQL extracts and AWS S3 ingestion feeding Power BI dashboards with refresh schedules

### Data Scientist – AI Engineering

May 2024 – Dec 2024

Inke (SafeSpace), Orlando, FL (Remote)

- Increased deployment reliability 30% by orchestrating model training & release via Airflow CI/CD on AWS EC2, serving 100K+ daily requests
- Improved API p95 latency 30% by profiling Django endpoints, optimizing containers, and tuning VPC networking
- Standardized retraining with MLflow Model Registry and feature versioning, reducing drift incidents and rollback time

### Data Scientist

Jan 2022 – Jun 2023

S L Kulkarni Cyril Graphics Pvt. Ltd., Mumbai, India (Remote)

- Cut unplanned downtime 20% by training predictive maintenance models on IoT sensors & logs and deploying anomaly alerts for preventive service
- Built Flask + SQL real-time dashboard for machine health, improving response speed and visibility across shifts

## Education

Kent State University, Kent, OH

Aug 2023 – May 2025

**Master of Science, Data Science** | GPA: 3.4/4.0 *Relevant coursework:* Artificial Intelligence (TensorFlow, PyTorch, Keras), Statistical Learning, Applied & Computational Statistics

## Technical Skills

- **Languages:** Python, R, SQL
- **Machine Learning & Deep Learning:** scikit-learn, TensorFlow, PyTorch, XGBoost, CNNs, HMMs, Bayesian Modeling, Time Series (ARIMA/SARIMAX), Cross-Validation, EDA
- **Generative AI & LLM Systems:** Multi-Agent Systems, Google ADK, Gemini, GPT, Llama, RAG Pipelines, LLM-powered Assistants
- **Data Engineering & MLOps:** Pandas, Dask, ETL Pipelines, Airflow, Docker, AWS (S3, EC2, Lambda), MLflow, DVC, CI/CD (GitHub Actions), Flask/Django, FastAPI
- **Frontend & APIs:** React, Next.js, TypeScript, Tailwind, REST API Development, Pydantic Validation
- **Healthcare Data:** HL7/FHIR, HIPAA-aware design, ICD-10, CPT, HCPCS, Synthea (Synthetic EHR data)
- **Visualization:** Power BI, Tableau, matplotlib, seaborn, Excel (Power Query, Pivot Tables)

## Projects

### EchoHealth – Safe Healthcare Voice AI Assistant

Oct 2025

- Built a voice-enabled AI assistant using Python, Flask, and OpenAI Whisper to process clinical queries with a MedlinePlus-backed RAG knowledge base for safe retrieval
- Implemented speech-to-text and clinical entity recognition aligned with HL7/FHIR schema awareness to support safe conversational triage and health information retrieval